

Homework 2

Classification Metrics

Group 2

3/10/2021

Contents

Assignment Overview	1
Task 1	2
Task 2	4
Functions	4
Task 3: Accuracy Function	4
Task 4: Classification Error Rate Function	5
Task 5: Precision Function	5
Task 6: Sensitivity Function	6
Task 7: Specificity Function	7
Task 8: F1 Score Function	7
Task 9	7
Task 10	8
Task 11	8
Task 12	8
Task 13	9

Group 2 members: *Alice Friedman, Diego Correa, Jagdish Chhabria, Orli Khaimova, Richard Zheng, Stephen Haslett.*

Assignment Overview

In this homework assignment, you will work through various classification metrics. You will be asked to create functions in R to carry out the various calculations. You will also investigate some functions in packages that will let you obtain the equivalent results. Finally, you will create graphical output that also can be used to evaluate the output of classification models, such as binary logistic regression.

The data set has three key columns we will use:

- **class:** the actual class for the observation.
- **scored.class:** the predicted class for the observation (based on a threshold of 0.5)
- **scored.probability:** the predicted probability of success for the observation

Task 1

Download the classification output data set.

```
data_raw <- read.csv("https://raw.githubusercontent.com/Jagdish16/CUNY_DATA_621/main/homework_2/classif
```

```
data_raw%>%head(50)
```

##	pregnant	glucose	diastolic	skinfold	insulin	bmi	pedigree	age	class
## 1	7	124	70	33	215	25.5	0.161	37	0
## 2	2	122	76	27	200	35.9	0.483	26	0
## 3	3	107	62	13	48	22.9	0.678	23	1
## 4	1	91	64	24	0	29.2	0.192	21	0
## 5	4	83	86	19	0	29.3	0.317	34	0
## 6	1	100	74	12	46	19.5	0.149	28	0
## 7	9	89	62	0	0	22.5	0.142	33	0
## 8	8	120	78	0	0	25.0	0.409	64	0
## 9	1	79	60	42	48	43.5	0.678	23	0
## 10	2	123	48	32	165	42.1	0.520	26	0
## 11	5	88	78	30	0	27.6	0.258	37	0
## 12	5	108	72	43	75	36.1	0.263	33	0
## 13	13	76	60	0	0	32.8	0.180	41	0
## 14	0	100	70	26	50	30.8	0.597	21	0
## 15	7	194	68	28	0	35.9	0.745	41	1
## 16	12	92	62	7	258	27.6	0.926	44	1
## 17	0	173	78	32	265	46.5	1.159	58	0
## 18	3	171	72	33	135	33.3	0.199	24	1
## 19	8	196	76	29	280	37.5	0.605	57	1
## 20	5	99	74	27	0	29.0	0.203	32	0
## 21	2	100	70	52	57	40.5	0.677	25	0
## 22	3	111	62	0	0	22.6	0.142	21	0
## 23	1	119	54	13	50	22.3	0.205	24	0
## 24	1	138	82	0	0	40.1	0.236	28	0
## 25	0	189	104	25	0	34.3	0.435	41	1
## 26	3	130	78	23	79	28.4	0.323	34	1
## 27	9	102	76	37	0	32.9	0.665	46	1
## 28	0	151	90	46	0	42.1	0.371	21	1
## 29	1	71	48	18	76	20.4	0.323	22	0
## 30	0	101	64	17	0	21.0	0.252	21	0
## 31	3	116	74	15	105	26.3	0.107	24	0
## 32	6	107	88	0	0	36.8	0.727	31	0
## 33	1	128	88	39	110	36.5	1.057	37	1
## 34	0	111	65	0	0	24.6	0.660	31	0
## 35	7	187	50	33	392	33.9	0.826	34	1
## 36	0	180	90	26	90	36.5	0.314	35	1
## 37	5	139	64	35	140	28.6	0.411	26	0
## 38	8	126	74	38	75	25.9	0.162	39	0
## 39	1	196	76	36	249	36.5	0.875	29	1
## 40	10	75	82	0	0	33.3	0.263	38	0
## 41	0	102	64	46	78	40.6	0.496	21	0
## 42	1	90	68	8	0	24.5	1.138	36	0
## 43	1	112	72	30	176	34.4	0.528	25	0
## 44	2	130	96	0	0	22.6	0.268	21	0

## 45	8	100	76	0	0	38.7	0.190	42	0							
## 46	3	89	74	16	85	30.4	0.551	38	0							
## 47	0	125	96	0	0	22.5	0.262	21	0							
## 48	2	91	62	0	0	27.3	0.525	22	0							
## 49	7	114	64	0	0	27.4	0.732	34	1							
## 50	7	136	90	0	0	29.9	0.210	50	0							
##	scored.class		scored.probability													
## 1	0	0.32845226														
## 2	0	0.27319044														
## 3	0	0.10966039														
## 4	0	0.05599835														
## 5	0	0.10049072														
## 6	0	0.05515460														
## 7	0	0.10711542														
## 8	0	0.45994744														
## 9	0	0.11702368														
## 10	0	0.31536320														
## 11	0	0.12518925														
## 12	0	0.27062482														
## 13	0	0.20980960														
## 14	0	0.09358589														
## 15	1	0.88484573														
## 16	0	0.39665216														
## 17	1	0.89139491														
## 18	1	0.53454900														
## 19	1	0.94633418														
## 20	0	0.14491618														
## 21	0	0.21763796														
## 22	0	0.07521357														
## 23	0	0.08843254														
## 24	0	0.30346820														
## 25	1	0.72448003														
## 26	0	0.27497369														
## 27	0	0.42486483														
## 28	0	0.43092552														
## 29	0	0.02322803														
## 30	0	0.04596084														
## 31	0	0.12798534														
## 32	0	0.29933706														
## 33	0	0.45909503														
## 34	0	0.10479581														
## 35	1	0.86309177														
## 36	1	0.63997495														
## 37	0	0.35818434														
## 38	0	0.37216467														
## 39	1	0.81110322														
## 40	0	0.16812736														
## 41	0	0.15127796														
## 42	0	0.10700703														
## 43	0	0.18796139														
## 44	0	0.13719711														
## 45	0	0.30047491														
## 46	0	0.13688715														
## 47	0	0.09786911														

```
## 48      0      0.06290701
## 49      0      0.26941931
## 50      0      0.48854279
```

Task 2

Use the `table()` function to get the raw confusion matrix for this scored dataset. Make sure you understand the output. In particular, do the rows represent the actual or predicted class? The columns?

```
# For the confusion matrix, we are only interested in the class and scored.class variables,
# so we select only these variables and ignore the rest.
confusion_matrix_table <- data_raw %>%
  select(class, scored.class)

# For readability purposes, rename 'scored.class' to Predicted, and 'class' to Actual.
dplyr::rename(confusion_matrix_table, Predicted = scored.class, Actual = class) %>%
  # Convert numeric boolean values to human readable values.
  mutate(Predicted = recode(Predicted,
                            '0' = 'Negative',
                            '1' = 'Positive'),
         Actual = recode(Actual,
                         '0' = 'Negative',
                         '1' = 'Positive')) %>%
  table()
```

```
##          Predicted
## Actual   Negative Positive
## Negative    119      5
## Positive    30     27
```

Functions

```
# We only need the class and scored.class variables from the dataset so we
# extract them and leave everything else.
data <- data_raw %>%
  select(class, scored.class)
```

Task 3: Accuracy Function

Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the accuracy of the predictions.

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + TN + FN}$$

```
accuracy <- function(df, col1, col2) {
  true = df[, col1]
  predict = df[, col2]
  # total events
  len = length(true)
```

```

# total correct predictions
correct = 0
for (i in seq(len)){
  if (true[i] == predict[i]){
    correct = correct + 1
  }
}
# accuracy
return (correct/len)
}
#example
accuracy(data_raw,'class','scored.class')

```

```
## [1] 0.8066298
```

Task 4: Classification Error Rate Function

Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the classification error rate of the predictions.

$$ClassificationErrorRate = \frac{FP + FN}{TP + FP + TN + FN}$$

```

class_error_rate <- function(df,col1,col2) {
  true = df[,col1]
  predict = df[,col2]
  # total events
  len = length(true)
  # total errors
  error = 0
  for (i in seq(len)){
    if (true[i] != predict[i]){
      error = error + 1
    }
  }
  # error rate
  return (error/len)
}
#example
class_error_rate(data_raw,'class','scored.class')

```

```
## [1] 0.1933702
```

Task 5: Precision Function

Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the precision of the predictions.

$$Precision = \frac{TP}{TP + FP}$$

```

#' Precision
#'
#' Given a dataset of actual and predicted classifications,
#' returns the precision of the predictions.
#'
#' @param data A dataset of actual and predicted classifications.
#'
#' @return Precision of predictions as a numeric value rounded to 2 decimal places.
precision <- function(data) {
  # Calculate the total number of true positives in the dataset.
  true_positive <- sum(data$class == 1 & data$scored.class == 1)

  # Calculate the total number of false positives in the dataset.
  false_positive <- sum(data$class == 0 & data$scored.class == 1)

  # Perform the precision calculation and round the result to 2 decimal places.
  prediction_precision <- round(true_positive / (true_positive + false_positive), 2)

  return(prediction_precision)
}

# Call the function to provide example output.
precision(data)

```

```
## [1] 0.84
```

Task 6: Sensitivity Function

Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the sensitivity of the predictions. Sensitivity is also known as recall.

$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

```

#note: there is a built in function in package caret called sensitivity
# head(data)

sensitivity <- function(data) {

  true_positive <- sum(data$class == 1 & data$scored.class == 1)
  false_negative <- sum(data$class == 1 & data$scored.class == 0)

  sensitivity <- true_positive / (true_positive + false_negative)

  return(sensitivity)
}

sensitivity(data)

```

```
## [1] 0.4736842
```

Task 7: Specificity Function

Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the specificity of the predictions.

$$\text{Specificity} = \frac{TN}{TN + FP}$$

```
specificity <- function(data){  
  
  true_negative <- sum(data$scored.class == 0 & data$class == 0)  
  false_positive <- sum(data$scored.class == 1 & data$class == 0)  
  
  specificity <- true_negative / (true_negative + false_positive)  
  
  return(specificity)  
}  
  
specificity(data)
```

```
## [1] 0.9596774
```

Task 8: F1 Score Function

Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the F1 score of the predictions.

$$F1Score = \frac{2 \times Precision \times Sensitivity}{Precision + Sensitivity}$$

```
# Should be based on the previous functions, so something like the below...  
f1_score<-function(data){  
  sens<-sensitivity(data)  
  prec<-precision(data)  
  f1<-2*sens*prec/(prec+sens)  
  return(f1)  
}  
  
f1_score(data)
```

```
## [1] 0.6057692
```

Task 9

What are the bounds on the F1 score? Show that the F1 score will always be between 0 and 1.

```
#data1<-data.frame(rep(0,5),rep(0,5))  
#data1<-data.frame(rep(0,5),rep(1,5))  
#data1<-data.frame(rep(1,5),rep(0,5))  
data1<-data.frame(rep(1,5),rep(1,5))  
colnames(data1)<-c('class','scored.class')  
data1
```

```
##   class scored.class
## 1     1           1
## 2     1           1
## 3     1           1
## 4     1           1
## 5     1           1
```

```
precision(data1)
```

```
## [1] 1
```

```
sensitivity(data1)
```

```
## [1] 1
```

```
f1_score(data1)
```

```
## [1] 1
```

Task 10

Write a function that generates an ROC curve from a data set with a true classification column (*class* in our example) and a probability column (*scored.probability* in our example). Your function should return a list that includes the plot of the ROC curve and a vector that contains the calculated area under the curve (AUC). Note that I recommend using a sequence of thresholds ranging from 0 to 1 at 0.01 intervals.

Task 11

Use your **created R functions** and the provided classification output data set to produce all of the classification metrics discussed above.

Task 12

Investigate the **caret** package. In particular, consider the functions *confusionMatrix*, *sensitivity*, and *specificity*. Apply the functions to the data set. How do the results compare with your own functions?

```
#convert the variables into factors as needed for the confusionMatrix
data <- data %>%
  mutate(scored.class = as.factor(scored.class),
         class = as.factor(class))

confusionMatrix(data$scored.class, data$class, positive = "1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 119  30
##           1   5  27
```



```
##
##           Accuracy : 0.8066
##           95% CI : (0.7415, 0.8615)
##      No Information Rate : 0.6851
##      P-Value [Acc > NIR] : 0.0001712
##
##           Kappa : 0.4916
##
##  McNemar's Test P-Value : 4.976e-05
##
##           Sensitivity : 0.4737
##           Specificity : 0.9597
##      Pos Pred Value : 0.8438
##      Neg Pred Value : 0.7987
##           Prevalence : 0.3149
##      Detection Rate : 0.1492
##      Detection Prevalence : 0.1768
##      Balanced Accuracy : 0.7167
##
##      'Positive' Class : 1
##
```

```
caret::sensitivity(data$score.class, data$class, positive = "1")
```

```
## [1] 0.4736842
```

```
caret::specificity(data$score.class, data$class, negative = "0")
```

```
## [1] 0.9596774
```

Task 13

Investigate the **pROC** package. Use it to generate an ROC curve for the data set. How do the results compare with your own functions?

```
pROC <- roc(data_raw$class, data_raw$score.probability)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(pROC, main = "pROC curve")
```

